

# Operating Systems

## C Programming Refresh

Jyrki Katajainen <jyrki@di.ku.dk>

### Promised by Jost:

pointers, malloc, free, essential data structures


### Actual contents:

motivation \* C machine \* pointers \* types \* kernel C \* Turing-computable functions \* syntactic sugar \* memory management \* pointers in action: stack \* macros \* inline assembly \* programming style \* idioms \* reflection

# Motivation

---

You should know C so well that you can understand the operating-system concepts explained in the book!




## C code

```
list_del(&p->list);  
kfree(p);
```

# Reading

---

- Brian W. Kernighan & Dennis M. Ritchie, *The C Programming Language* (1988)



## Optional

- Brian W. Kernighan & Rob Pike, *The Practice of Programming* (1999)
- Dennis M. Ritchie, The development of the C language, *ACM SIGPLAN Notices* (1993)
- S. Sandeep, *GCC-Inline-Assembly-HOWTO* (2003)

# Disclaimer

---


I will concentrate on the fundamental issues, not any specific details. You have to read the textbook and practise to learn the language.

E.g. I will skip the following topics:

- bit manipulation [K&R, Section 2.9]
- command-line arguments [K&R, Section 5.10]
- structures [K&R, Ch. 6]
- input and output [K&R, Ch. 7]
- Unix interface [K&R, Ch. 8]
- library resources [K&R, App. B]

# Knowledge levels

---



# Skills you need

---

- Read old code
- Read badly-written code
- Read code written in a strange way
- Read low-level assembly code

**Guru:** Replace “Read” with “Debug and correct”


```
jyrki@Janus$ cat 100.c
_( __, ___, ___ ) { ___ / __ <= 1 ? _ ( __, ___ + 1, ___ ) : ! ( ___ % __ ) ? _ ( __, ___ + 1, 0 ) : ___ % __
== ___ /
__ && ! ___ ? ( printf ( "%d\t", ___ / __ ), _ ( __, ___ + 1, 0 ) ) : ___ % __ > 1 && ___ % __ < ___ / __ ?
_ ( __, 1 +
___, ___ + ! ( ___ / __ % ( ___ % __ ) ) ) : ___ < __ * __ ? _ ( __, ___ + 1, ___ ) : 0 ; } main () { _ ( 100,
0, 0 ) ; }
jyrki@Janus$ gcc 100.c
100.c: In function ‘_’:
100.c:3:12: warning: incompatible implicit declaration of function ‘printf’
__ && ! ___ ? ( printf ( "%d\t", ___ / __ ), _ ( __, ___ + 1, 0 ) ) : ___ % __ > 1 && ___ % __ < ___ / __ ?
^
jyrki@Janus$ ./a.out
2 3 5 7 11 13 17 19 23 29 31 37 41 43 47 53 59 61 67 71 73 79 83 89
97
```

# C machine

---



## Memory



# Pointers

---

A **pointer** is a variable that contains the address of a variable.


**Address:** Universitetsparken 5

**Building at that address:** HCØ

**p:** 0x7fff30788da4

**\*p:** 5

0x7fff30788da4



```
int x = 5;  
int* p = &x;  
assert(*p == 5);
```



# Types

---

In C, data declarations are read from right to left (and from the inside out). Keep this in mind when writing data declarations; typedefs can be used to improve readability.

`int const* p;` Pointer to a constant integer, i.e. the integer pointed to does not change.

`int* const q;` Constant pointer to an integer, i.e. the pointer is kept constant but the underlying integer can be modified.

```
typedef char* routine(char*, char const*);  
routine* copy = copy_string;
```

# Online exercise

---

What can you say about the following function signature?

```
#include <stdint.h>
...
TID_t thread_create(void (*func)(uint32_t), uint32_t arg)
```

[Buenos roadmap 2012, p. 19]

**Warning:** Your compiler understands C99 or C11, whereas our textbook describes ANSI C completed in 1989.

# Expressions

---

- Expressions combine variables and constants to produce new values.
  - Every expression has a value.
  - The value of a relational or logical expression is 1 (true) or 0 (false)
  - 0 means false and any other value means true.
  - Be aware of the evaluation order of operators in an expression!
- `x = 0` is an expression
  - `x = 0;` is a statement with the terminator ;
  - `x = 0` has the value 0 (the value of the left-hand side after the assignment)

# Kernel of the C language

---

A *kernel-C program* is a sequence of statements that are executed sequentially unless the order is altered by a `while` statement.

The kernel of the C language has

- assignments,
- while statements, and
- nothing else.

Observe that a sequence  $\{\text{statement}_1 \text{ statement}_2 \dots\}$  is also a statement.

$x$ : variable;  $y, z$ : variable | constant

$p$ : pointer variable

$\mathcal{A}$ :  $\{+, -, *, /, \%\}$  (arithmetic)

$\mathcal{B}$ :  $\{\&, |, ^, \ll, \gg\}$  (bitwise)

$\mathcal{C}$ :  $\{<, <=, ==, !=, >, >=\}$  (comparison)

$\mathcal{U} = \{-, \sim, \&\}$  (unary)

**Load:**  $x = *p;$

**Store:**  $*p = y;$

**Move:**  $x = y;$

**Unary operation:**  $\ominus \in \mathcal{U}$

$x = \ominus y;$

**Binary operation:**  $\oplus \in \mathcal{A} \cup \mathcal{B} \cup \mathcal{C}$

$x = y \oplus z;$

**While loop:**  $\triangleleft \in \mathcal{C}$

while ( $y \triangleleft z$ )

statement

# Theorem

---

$n, m$ : positive integers

$\mathbb{B}$ :  $\{0, 1\}$

task: compute a function mapping  $\mathbb{B}^n$  to  $\mathbb{B}^m$

Kernel C can be used to perform exactly the same tasks as your favourite programming language X++.

# Fundamental theorem

---

$P$ : some program

$X$ : task performed by  $P$

$\kappa$ : # kernel-C statements in  $P$

$T(n)$ : running time of  $P$  for an input of size  $n$

There exists a program  $P'$

1. that has only one `while` statement,
2. that performs  $X$ ,
3. whose length is  $O(\kappa)$ ,
4. whose running time is  $O(\kappa T(n))$ .

For a proof, see, e.g. [Elmasry & Katajainen 2013]

# Syntactic sugar

---

```
if (expression)
    statement1
else
    statement2
```

```
done = 0;
if (expression)
    statement1
done = 1;
if (! done && ! expression)
    statement2
```

```
if (expression)
    statement
```


```
done = 0;
while (! done && expression)
    statement
done = 1
```

```
for (expression1; expression2; expression3)
    statement
```

```
expression1;
while (expression2)
    statement
expression3;
```

# Advanced features

---





# Memory management

---

“Careful memory management—which includes releasing memory to prevent **memory leaks**—is crucial when developing kernel-level code.”

```
/* allocate n bytes from the heap */  
void* malloc(unsigned int n);  
  
/* release the block pointed to by p */  
void free(void* p);
```

For an implementation of a storage allocator, see [K & R, Section 8.7]

# Kernel data structures

---

We expect that you know the following data structures (and know how to program them in C):

- singly-linked lists
- doubly-linked lists (Assignment 1)
- stack
- queue
- red-black tree
- hash tables
- bit vectors.

For a recap, read [Our textbook, Section 2.10]

# Macros

---

Can be used

- to define constants,
- for conditional compilation, and
- to define macro functions.

```
#ifndef __bool_true_false_are_defined || __bool_true_false_are_defined == 0
    typedef int bool;
    #define false 0
    #define true !(false)
#endif
```

# Inline assembly (gcc specific)

---

## Basic inline assembly

```
asm("assembly code");
```

### Example

```
asm("movl %ecx %eax");
```

### Warning

The AT&T syntax used all over; information flows from left to right. In the above example the contents of ecx is moved to eax.

## Extended inline assembly

```
asm( assembler template
    : output operands /* optional */
    : input operands /* optional */
    : list of clobbered /* optional */
);
```

### Example

```
int a = 10;
int b;
asm(
    "movl %1, %%eax\n"
    "movl %%eax, %0"
    : "=r" (b) /* output */
    : "r" (a) /* input */
    : "%eax" /* clobbered */
);
```

# Why is good programming style important?

---

**Clarity:** Well-written code is easier to read and to understand.

**Correctness:** Sloppy code is often broken.

**Simplicity:** Well-written code is likely to be smaller than code that has been carelessly tossed together and never polished.

**Maintainability:** The best layout schemes hold up well under code modifications.

[Kernighan & Pike 1999, §1]

# Idiomatic code

---

- Use expressions that are natural for a C programmer.
- Establish conventions in non-critical areas so that you can focus your creative energies in the places that count.
- Methods and tools that emphasize human discipline are especially effective. Form is liberating!

# Document-comment idiom

---

Write a brief comment (if any) before an important function giving a high-level description what the function does, not how it does it.

```
/* count lines, words, and characters in input */  
int main() {  
    ...  
    return 0;  
}
```

# Variable-naming conventions

---

- `c` is a character variable
- `i`, `j`, and `k` are integer indexes
- `n` is a number of something
- `p` and `q` are pointers
- `s` is a string
- Preprocessor macros are all in upper case as in `ALL_CAPS`
- Variable and routine names are all in lower case as in `all_lower_case`
- The underscore (`_`) character is used as a separator between words as above.



# Variable-naming idioms

---

- Use descriptive names for global variables, and short names for local variables.
- Select sensible names for your variables. A misleading name can result in mysterious bugs.
- Use active names for functions. Also, functions that return a Boolean should be named so that the return value is unambiguous (e.g. `is_empty`).

# Indentation idiom

---

Indentation should be idiomatic, too. Use 2-4 space indentation to show the logical structure of a program. If you use tabs, make sure that your code-beautifier and other tools handle them appropriately.


```
    for (++n; n != 100; ++n) {  
        field[n] = '\0';  
    }
```

```
? for (  
?     n++;  
?     n != 100;  
?     field[n++] = '\0';  
? )  
? {  
?     ;  
? }
```

# Which brace style is best?


---

## Pure blocks:



```
while (colour == red) {
    statement1
    statement2
    ...
}
```

## Braces as block boundaries:



```
while (colour == red)
{
    statement1
    statement2
    ...
}
```

# Basic-loop idiom

---

Most loops, where the exit test is performed at the beginning, should have the form:

```
for (i = 0; i != n; ++i) {  
    ...  
}
```

```
for (p = list; p != 0; p = (*p).next) {  
    ...  
}
```

```
? i = 0;  
? while (i <= n - 1) {  
?     ...  
? }
```

```
? for (i = 0; i < n; ) {  
?     ...  
?     i++;  
? }
```

```
? for (i = n; --i >= 0; ) {  
?     ...  
? }
```

# Infinite-loop idiom

---

For infinite loops use one of the following alternatives; do not use other forms.

```
#include <stdbool.h>
...
while (true) {
    ...
}

for (;;) {
    ...
}

? while (1) {
?     ...
? }
```

# End-exit idiom

---

A do-while loop always executes at least once; in many cases that behaviour is a bug waiting to bite, but when it is needed, write `} in front of while to indicate clearly that while does not start a while loop.`

```
safety_counter = 0;
do {
    current = (*current).next;
    ...
    ++safety_counter;
    if (safety_counter >= SAFETY_LIMIT) {
        fprintf(stderr, "Internal error: Safety-counter violation\n");
        assert(false);
    }
} while ((*current).next != 0);
```

# Middle-exit idiom

---

Few programming languages provide direct support for a loop construct whose exit point is in the middle. An idiomatic way of writing such a loop is shown in the following example:

```
while (true) {  
    c = getchar();  
    if (c == EOF) {  
        break;  
    }  
    putchar();  
}
```

```
? while ((c = getchar()) != EOF) {  
?     putchar();  
? }
```

# Switch-break idiom

---

Cases should almost always end with a break, with the rare exceptions commented. For an unusual structure, a sequence of else-if statements can be even clearer.

```
if (c == '-') {
    sign = -1;
    c = getchar();
}
else if (c == '+') {
    c = getchar();
}
else if (c != '.' && !isdigit(c)) {
    return 0;
}
```

```
? switch (c) {
? case '-': sign = -1;
? case '+': c = getchar();
? case '.': break;
? default: if (!isdigit(c))
?             return 0;
? }
```

```
? switch (c) {
? case '-':
?     sign = -1;
?     /* fall through */
? case '+':
?     c = getchar();
?     break;
? case '.':
?     break;
? default:
?     if (!isdigit(c)) {
?         return 0;
?     }
?     break;
? }
```



# Operation-per-line idiom

---

Avoid using multiple operations per line; in particular, be careful with statements that have side effects.

```
char* copy_string(char* s, char const* t) {  
    char* const r = s;  
    while (*t != '\0') {  
        *s = *t;  
        ++s;  
        ++t;  
    }  
    *s = '\0';  
    return r;  
}
```

```
? void strcpy(char *s, char *t) {  
?   while (*s++ = *t++)  
?       ;  
? }
```

# Declaration-per-line idiom

---

Use only one data declaration per line.

```
FILE* input_file;  
FILE* output_file;
```

```
? FILE *input_file, *output_file;
```

Separate the type names and variable names clearly in data declarations.

# Memory-allocation idiom

---

In a real program, the return value of `malloc`, `realloc`, or any other allocation routine should always be checked.

```
p = (char*) malloc(strlen(buffer) + 1);
if (p == 0) {
    reset();
    return 0;
}
strcpy(p, buffer);

? p = malloc(strlen(buffer) + 1);
? strcpy(p, buffer);
```

# Some comments on readability

---

- Operators:** Use spaces around operators: `x = y + z;`
- ,:** Add a single space after commas: `a = f(x, y, z);`
- !:** Add a single space after `!`; negations are hard to understand and should be avoided.
- if:** Deeply nested `if` statements are difficult to follow.
- >:** Arrows clutter the code; my preference is `(*p).next`.

Instead of following the rules slavishly, it is more important to be consistent.

# Some comments on maintainability

---

- Give names to magic numbers.
- Be prepared for changes.
- Use a commenting style that is easy to maintain.

```
#define SIZE(array) (sizeof(array) / sizeof(array[0]))

char buffer[100];

for (i = 0; i != SIZE(buffer); ++i) {
    // Use braces even if the loop body only contains one statement.
}
```

# Conclusion

---

C is quirky, flawed, and an enormous success.

[Ritchie 1993]

# Critique of C

---

- Programmers have to do their own memory management—to declare variables, to manage pointer-chained lists, to dimension buffers, to detect and prevent buffer overruns, and to allocate and deallocate dynamic storage
- Little support for modularization
- No support for generic functions

- No string data type
- No support for variable-sized arrays

- Array syntax is a living fossil

```
int f(int a[]);  
/* type of a is int* */
```

- Other syntax peculiarities

```
int* (*g)(int);  
/* pointer to function ↔  
   returning int* */
```

- Semantic peculiarities

```
r = a >> 4 + 5;  
/* shift right by 9 */
```

- Difficult to write numerical libraries.

# Whence success

---

- Integration with the UNIX operating system
- A small and simple language
- Translatable with small and simple compilers
- Its types and operations are well-grounded in those provided by real machines
- Sufficiently abstract from machine details to achieve program portability
- Sufficiently expressive for describing time- and space-efficient programs
- Has remained remarkably stable
- Today, useful as a portable high-level assembly language.

Adopted from [Ritchie 1993]